

AI Deep Learning: Convolutional Neural Networks (II)


Deep Learning with Big Data

Thuan L Nguyen, PhD

AI Deep Learning: Convolutional Neural Networks (CNN)

1. AI Deep Learning: Convolutional Neural Networks Overview
2. AI Deep Learning: Properties of Images
3. AI Deep Learning: Convolution Layers in Convolutional Neural Networks
4. AI Deep Learning: Pooling Layers in Convolutional Neural Networks
5. AI Deep Learning: Convolutional Neural Networks: Architectures: Overview
6. AI Deep Learning: Convolutional Neural Networks: Architectures: Layers
7. AI Deep Learning: Build and Train Convolutional Neural Networks

AI Deep Learning: Convolutional Neural Networks (CNN)



AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Convolution(al) Operation, Feature Maps, and Receptive Fields

- The **convolution operation** is performed by sliding the filter over the input.
- At every location, **element-wise matrix multiplication** is done and the **results** are **summed**.
- The sum goes into the **feature map**.
- The green area where the convolution operation takes place is called the **receptive field**.
 - Due to the size of the filter, the receptive field is also 3x3.

1x1	1x0	1x1	0	0
0x0	1x1	1x0	1	0
0x1	0x0	1x1	1	1
0	0	1	1	0
0	1	1	0	0

Input x Filter

4		

Feature Map

AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Examples of Convolution(al) Operation

- This was an example convolution operation shown in 2D using a 3x3 filter:

Original image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Convolutional filter 1

1	0	1
0	1	0
1	0	1

Convolving the image

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Result

4		

AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Examples of Convolution(al) Operation

- This was an example convolution operation shown in 2D using a 3x3 filter:

Original image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Convolutional filter 1

1	0	1
0	1	0
1	0	1

Convolving the image

1	1 _{x1}	1 _{x0}	0 _{x1}	0
0	1 _{x0}	1 _{x1}	1 _{x0}	0
0	0 _{x1}	1 _{x0}	1 _{x1}	1
0	0	1	1	0
0	1	1	0	0

Result

4	3	

AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Examples of Convolution(al) Operation

- This was an example convolution operation shown in 2D using a 3x3 filter:

Original image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Convolutional filter 1

1	0	1
0	1	0
1	0	1

Convolving the image

1	1	1	0	0
0	1	1	1	0
0	0	1 _{x1}	1 _{x0}	1 _{x1}
0	0	1 _{x0}	1 _{x1}	0 _{x0}
0	1	1 _{x1}	0 _{x0}	0 _{x1}

Result

4	3	4
2	4	3
2	3	4

AI Deep Learning: Convolutional Neural Networks (CNN)


CNN: Convolution(al) Operations & Final Outputs

- In previous slides is an example convolution operation shown in 2D using a 3x3 filter.
- **In reality:**
 - An image is represented as a **3D matrix** with dimensions of **height, width and depth**,
 - Where **depth** corresponds to **color channels (RGB)**.
- A 2D **convolutional layer** → **2D filter** has a specific **height and width**, like **3x3** or **5x5**,
- With actual convolution operations:
 - **Multiple convolutions** are performed on an input:
 - **Each** using a **different filter** and resulting in **a distinct feature map**.
- We then **stack** all these feature maps together.
 - To create the **final output** of the **convolution layer**.

AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Convolution(al) Operations & Final Outputs


- Convolution operations result in a set of feature maps.
- All these feature maps are stacked together to create the **final output** of the **convolution layer**.




AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Convolution(al) Operations: Stride


- In convolution operations, **stride** specifies **how much** we move the convolution filter **at each step**.
- By **default**, the value is **1**, as you can see in the figure below.




Stride 1



Feature Map



Stride 1




Feature Map


AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Convolution(al) Operations: Stride


- It is possible to use **bigger strides**
 - That makes the resulting feature map **smaller**
 - It's likely to skip over potential locations.
 - That results in less overlap between the receptive fields.
- The following figure demonstrates a stride of 2.




Stride 2



Feature Map




Stride 2




Feature Map

AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Convolution(al) Operations: Padding



Stride 2



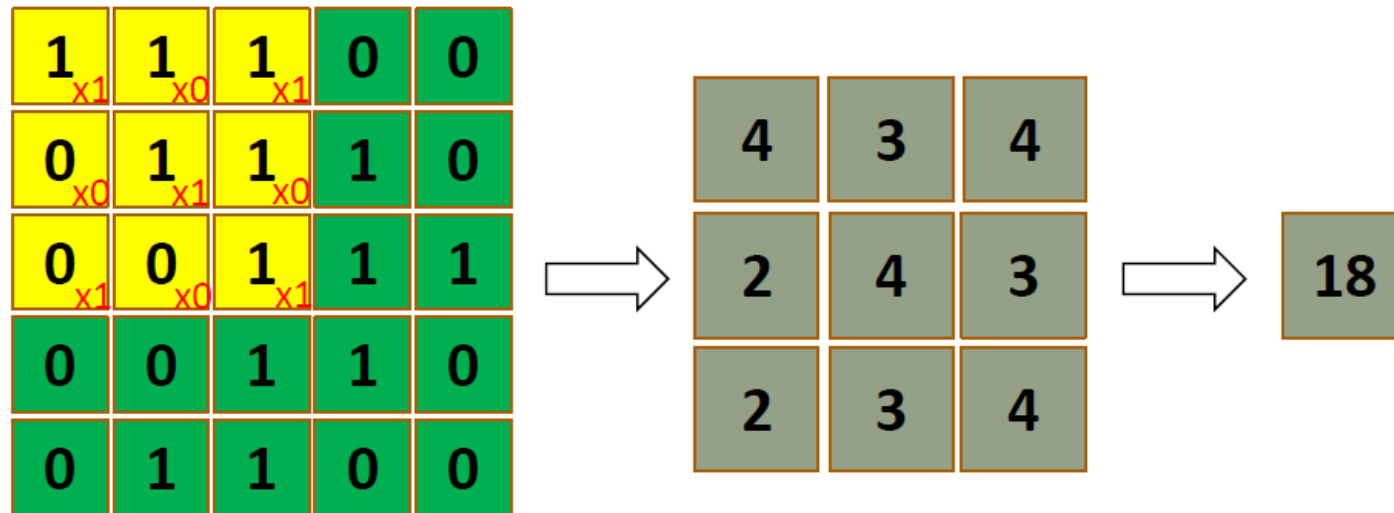
Feature Map

- So, the **size** of the feature map is **smaller** than the input:
 - Because the **convolution filter** needs to be contained in the input.
- It's GOOD!
 - Convolution operations reduce dimensionality.
 - Fewer dimensions, simpler & faster computation.

AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Convolution(al) Operations: Padding


- BUT!
 - It's also a big problem. What?
 - The image inputs get smaller and smaller
 - → Details are lost → recognition accuracy drops
 - → Run out of “latent pixels”
 - → Convolutional neural networks: not too deep architectures




- If we want to **maintain** the same dimensionality:
 - Use **padding** to surround the input with zeros.

AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Convolution(al) Operations: Padding



Stride 1 with Padding



Feature Map

- The gray area around the input is the **padding**.
- We either **pad with zeros or the values on the edge**.
 - Padding makes the dimensionality of the feature map **match** the input.
- **Padding** is commonly used in **convolutional neural networks** to **preserve the size of the feature maps**.
 - Otherwise, the size of feature maps would shrink at each layer, which is not desirable.

AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Convolution(al) Operations: Zero Padding

- HOWTO calculate how many layers of values for padding?
 - Given **h_filter**: height of the filter; **w_filter**: width of the filter
 - Number of layers of values for padding = $(h_filter - 1) / 2$ OR $(w_filter - 1) / 2$
 - For example, with $h_filter = 3$ and $w_filter = 3$, layer of padding = 1

0	0	0	0	0	0	0
0	1	1	1	0	0	0
0	0	1	1	1	0	0
0	0	0	1	1	1	0
0	0	0	1	1	0	0
0	0	1	1	0	0	0
0	0	0	0	0	0	0

*

0	0	1
0	1	1
1	1	1

=

1	1	2	0	0
0	1	1	1	0
0	0	1	2	1
1	0	2	1	0
0	1	1	3	0

AI Deep Learning: Convolutional Neural Networks (CNN)


CNN: Convolution(al) Operations: Suggested Best Practices

- Image size:
 - Resize the image to have a size in the **power of 2**, e.g., 32x32, 256x256
- Convolution filter:
 - A filter of (h, w) = **3 x 3** should work well with deep neural networks
- Stride:
 - Use the default stride: **s = 1**
- Padding:
 - Add **1 layer** of zero padding

AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Convolution(al) Operations: An Example of 3D Filters

- First, to be simple, a convolution using a single filter.





- Given a 32x32x3 image and a filter of size 5x5x3
 - NOTES: *the depth of the convolution filter matches the depth of the image, both being 3.*
- When the filter is at a particular location, it covers a small volume of the input, and the convolution operation is performed as above.
- In this convolution operation, the sum of matrix multiply in 3D – instead of 2D – is done; the result is still a scalar.
- Slide the filter over the input as above and perform the convolution at every location aggregating the result in a feature map.
- This feature map is of size 32x32x1, shown as the red slice on the right.

AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Convolution(al) Operations: An Example of 3D Filters

- If we used 10 different filters we would have 10 feature maps of size $32 \times 32 \times 1$:
 - Stacking them along the depth dimension produces the final output of the convolution layer:
 - a volume of size $32 \times 32 \times 10$, shown as the large blue box on the right.







AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Convolution(al) Operations: An Example of 3D Filters

- NOTES:
 - *The height and width of the feature map are unchanged and still 32 due to padding.*
- To help with visualization:
 - Slide the filter over the input as shown in the simple scenario of the 2D filter.
 - At each location, a scalar is calculated.
 - The set of all these scalars make the feature map.
- The figures show the sliding operation at 4 locations.
 - However, in reality, the convolution operation is performed over the entire input, i.e., all over the blue rectangle to the left.






AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Convolution(al) Operations: An Example of 3D Filters

- Let' see how two feature maps are stacked along the depth dimension:
 - The convolution operation for each filter is performed independently.
 - The resulting feature maps are separated.



AI Deep Learning: Convolutional Neural Networks (CNN)


CNN: Pooling

- After a convolution operation we usually perform **pooling** to reduce the dimensionality.
 - This enables us to reduce the number of parameters
 - That helps both **shorten the training time** and combat **overfitting**.
- **Pooling layers** down-sample each feature map independently.
 - That **reduces the height and width, keeping** the depth intact.
- The most **common** type of pooling is **max pooling**.
 - That just takes the **max value** in the pooling window.
- **Different from** the convolution operation, **pooling has no parameters**.
 - It slides a window over its input, and simply takes the max or the average value in the window.
 - Similar to a convolution, we **specify the window size and stride**.

AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Pooling: Why?

- Pooling **aggregates** multiple values into a single value.
- Pooling: **Invariance** to small transformations
- Pooling **reduces the size** of the layer output/input to next layer.
 - So, faster computations
- Pooling **keeps most important information** for the next layer



AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Pooling: Max Pooling vs. Average Pooling

- Max pooling:
 - Keep the max value in the window
- Average pooling:
 - Calculate the average of all the values in the window and keep it.

1	4	3	6
2	1	0	9
2	2	7	7
5	3	3	6

4	9
5	7


1	4	1	6
2	3	0	9
1	2	7	1
4	1	0	2

2.5	4
2	2.5

AI Deep Learning: Convolutional Neural Networks (CNN)


CNN: Pooling: An Example with Max Pooling

- Here is the result of max pooling using a 2x2 window and stride 2
 - Each color denotes a different window.
 - Since both the window size and stride are 2, the windows are not overlapping.
- This window and stride configuration **halves the size of the feature map**.
 - This is the **main use case of pooling**, down-sampling the feature map while keeping the important information.



AI Deep Learning: Convolutional Neural Networks (CNN)


CNN: Pooling



- Consider the **feature map dimensions** **before** and **after** pooling:
 - If the input to the pooling layer has the dimensionality $32 \times 32 \times 10$, using the same pooling parameters described above, the result will be a $16 \times 16 \times 10$ feature map.
 - Both the **height** and **width** of the feature map are **halved**, but **the depth is not changed** because **pooling works independently** on each depth slice the input.

AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Pooling



- By **halving** the height and the width, we **reduced** the number of weights to **1/4** of the input.
 - Consider that to train a deep learning model, it is typical to deal with **millions of weights** in CNN architectures, this reduction is **a big deal**.
-
- In **convolutional neural network architectures**:
 - Convolution is done with **3x3 windows**, **stride 1** and with **padding**.
 - Pooling is typically performed with **2x2 windows**, **stride 2** and **no padding**.

AI Deep Learning: Convolutional Neural Networks (CNN)

CNN: Pooling

- Here is the result of max pooling using a 2x2 window and stride 2
 - Each color denotes a different window.
 - Since both the window size and stride are 2, the windows are not overlapping.
- This window and stride configuration **halves the size of the feature map**.
 - This is the **main use case of pooling**, **down-sampling the feature map** while **keeping the important information**.

